

# REDIRE — BASELINE REPORT

OLIVER RICHARDSON

MAKS CEGIELSKI-JOHNSON

## 1. TASK

Our interest in this problem is two-fold; we are interested in studying both recognition and generation of paraphrases. Hence we will need to discuss our current approaches to both of these problems. 4 The paraphrase recognition task is defined as attempting to determine whether two sentences are paraphrases of each other, returning a YES/NO answer.

The paraphrase generation task is defined as being given an input sentence, and trying to generate a paraphrase from that input.

## 2. ARCHITECTURE

Our current implementation (SIM) is inspired by the INIT method in a paper by Malakasiotis[1]. For the baseline, we used the Microsoft Research Paraphrase Corpus[3] (MSRP) The general structure of this implementation relies on transforming a pair of sentences to different linguistic representations and then computing a variety of similarity metrics between each representation pair. This procedure is described below.

- (1) First, read all sentence pairs  $(s_1, s_2)$  from a file.
- (2) Restructure each pair of sentences  $s_1$  and  $s_2$  into 10 new string transformations  $t_{i,1}$  and  $t_{i,2}$ , maintaining the original syntactic order of the sentence.
  - (a) String of the tokens
  - (b) String of the stems.
  - (c) String of the POS tags.
  - (d) String of the soundex codes[4].
  - (e) String of all the noun tokens
  - (f) String of all the noun stems
  - (g) String of all the noun soundex codes.
  - (h) String of all the verb tokens.
  - (i) String of all the verb stems.
  - (j) String of all the verb soundex codes.
- (3) For each string pair  $s_1$  and  $s_2$ , and for each transformation of that string  $t_{i,1}$  and  $t_{i,2}$ , we want to compute the similarity[2]  $d_i(t_{i,1}, t_{i,2})$ , using the metrics
  - (a) Levenshtein (word edit)
  - (b) Jaro-Winkler
  - (c) Manhattan & Euclidean
  - (d) Cosine
  - (e)  $n$ -gram ( $n = 3$ )
  - (f) Overlap
  - (g) Dice coefficient
  - (h) Jaccard coefficient
- (4) Computing the similarities between each transformation of the two strings results in 90 features.
- (5) Next, for each pair of sentences  $s_1$  and  $s_2$ , if  $s_1$  is longer than  $s_2$  we compute the substrings of  $s_1$  which have the same length as  $s_2$ . For each substring  $s'_1$  generated, we compute all of the similarities between  $s'_1$  and  $s_2$ . We try to find which substring results in the max

average similarity over all the similarity metrics, namely  $s_1^*$ . Then for  $s_1^*$  we compute each of the similarities between it and  $s_2$ , as well as the average of all of them, resulting in 10 features. We do this for the string transformations which result in tokens, stems, POS tags, and soundex codes. This results in an additional 40 features per sentence pair (total of 130 features so far). Note that we can just swap the two strings in the discussion above if  $s_2$  is larger than  $s_1$ .

- (6) Finally, we add three additional features to the current 130. We add a boolean feature for whether  $s_1$  contains a negation, and a feature for whether  $s_2$  contains negation. Then finally, we add a feature for the length ratio, defined as  $\frac{\min(|s_1|, |s_2|)}{\max(|s_1|, |s_2|)}$ .
- (7) Malakasiotis then discusses using a dependency grammar for features (DEP), which we found easy to implement as well. We can compute the dependency parses for both sentence  $s_1$  and sentence  $s_2$ . Using the dependency parses, we can calculate the following

$$R_1 = \frac{|\text{common dependencies}|}{|\text{dependencies in } s_1|} \quad R_2 = \frac{|\text{common dependencies}|}{|\text{dependencies in } s_2|}$$

where  $R_1$  and  $R_2$  can both be used as features. We can then also use the harmonic mean  $\frac{2R_1R_2}{R_1+R_2}$  as the final feature.

### 3. EXPERIMENTS

Having defined how to get our similarity-based features, we can start experimenting with different combinations. First, let us define a baseline experiment, (BASE), for which we will use the Levenshtein distance and a simple linear classifier. This baseline will allow us to see how a naïve approach will work on the our dataset.

SIM is the method discussed above which uses string transformations and string similarity metrics. This experiment will allow us to see how well using only string similarity will work for paraphrasing. DEP introduces the dependency grammar for a few features, allowing the use of the structure of the sentence for paraphrase recognition.

### 4. RESULTS

Method	Accuracy	Precision	Recall	F-Score
BASE	0.64811	0.68442	0.87358	0.76752
SIM	0.66492	0.72814	0.79163	0.75856
SIM + DEP	0.66492	0.73224	0.78204	0.75632

TABLE 1. REDIRE results

Method	Accuracy	Precision	Recall	F-Score
BASE	0.6904	0.7242	0.8631	0.7876
INIT	0.7519	0.7851	0.8631	0.8223
INIT + WN	0.7548	0.7891	0.8614	0.8237
INIT + WN + DEP	0.7935	0.7891	0.8675	0.8288

TABLE 2. Malakasiotis results [1]

### 5. DISCUSSION

Table 1 contains all the results from our experiments. Table 2 contains all the results from the Malakasiotis paper, allowing us to compare how well our implementation did. One notable difference

is that our implementation does not use WordNet, while Malakasiotis does. This must be considered when comparing our SIM + DEP to Malakasiotis' INIT + WN + DEP.

Clearly, following the Malakasiotis paper has not given us comparable results, netting a 7% lower F-score using just similarity metrics. The next step would be to determine what is causing the much lower score. However, we are still happy with the progress we currently have. Clearly from Table 2 we can see that adding WordNet (WN) did not boost performance significantly, so we don't necessarily need to focus our attention on implementing this functionality.

## 6. CONCLUSIONS

Since both REDIRE and Malakasiotis only use naïve word similarity metrics for classification, we are interested in attempting to explore deeper means for paraphrase recognition. Using only similarity means that two sentences which are paraphrases of each other but don't share similar words will not be recognized as paraphrases. We want something that can relate the similarity between two paraphrases other than just the words.

One way we can attempt to achieve this is through the use of a dependency grammar and WordNet. Seeing as this is what the Malakasiotis paper does as well, we need to consider different ways of representing similarity using both tools. If we can achieve a adequate means of doing so, then not only will this help us boost performance for our paraphrase recognition, but will also allow us to begin to approach the problem of paraphrase generation.

## 7. CONTRIBUTIONS

TEXT 1

TEXT 2

## REFERENCES

- [1] Prodromos Malakasiotis. Paraphrase recognition using machine learning to combine similarity measures. In *Proceedings of the ACL-IJCNLP 2009 Student Research Workshop*, pages 27–35. Association for Computational Linguistics, 2009.
- [2] Prodromos Malakasiotis and Ion Androutsopoulos. Learning textual entailment using svms and string similarity measures. In *Proceedings of the ACL-PASCAL Workshop on Textual Entailment and Paraphrasing*, pages 42–47. Association for Computational Linguistics, 2007.
- [3] Microsoft Research. Microsoft research paraphrase corpus. <http://research.microsoft.com/en-us/downloads/607d14d9-20cd-47e3-85bc-a2f65cd28042/>.
- [4] R. Sedgewick and K. Wayne. Introduction to Programming in Java: Soundex. <http://introcs.cs.princeton.edu/java/31datatype/Soundex.java>.